

## Context

Jamboree has helped thousands of students make it to top colleges abroad. Be it GMAT, GRE or SAT, their unique problem-solving methods ensure maximum scores with minimum effort.

They recently launched a feature where students/learners can come to their website and check their probability of getting into the IVY league college. This feature estimates the chances of graduate admission from an Indian perspective.

## Problem Statement :

Help Jamboree in understanding what factors are important in graduate admissions and how these factors are interrelated among themselves. It will also help predict one's chances of admission given the rest of the variables.

## Column Profiling:

Serial No. (Unique row ID)

GRE Scores (out of 340)

TOEFL Scores (out of 120)

University Rating (out of 5)

Statement of Purpose and Letter of Recommendation Strength (out of 5)

Undergraduate GPA (out of 10)

Research Experience (either 0 or 1)

Chance of Admit (ranging from 0 to 1)

## Exploratory Data Analysis

## Linear Regression

In [67]:

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 %matplotlib inline
6 from matplotlib import figure
7 import warnings
8 warnings.filterwarnings("ignore")
9 import statsmodels.api as sm
10
11
```

In [68]:

```
1 data = pd.read_csv("Admission_Predict_Ver1.1.csv")
2 data.head()
```

Out[68]:

Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit	
0	1	337	118	4	4.5	4.5	9.65	1	0.92
1	2	324	107	4	4.0	4.5	8.87	1	0.76
2	3	316	104	3	3.0	3.5	8.00	1	0.72
3	4	322	110	3	3.5	2.5	8.67	1	0.80
4	5	314	103	2	2.0	3.0	8.21	0	0.65

In [69]:

```
1 data.sample(5)
```

Out[69]:

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
364	365	313	102	3	3.5	4.0	8.90	1	0.77
13	14	307	109	3	4.0	3.0	8.00	1	0.62
265	266	313	102	3	2.5	2.5	8.68	0	0.71
236	237	325	112	4	4.0	4.5	9.17	1	0.85
499	500	327	113	4	4.5	4.5	9.04	0	0.84

In [70]:

```
1 # Checking the shape of data
2 data.shape
```

Out[70]:

(500, 9)

In [71]:

```
1 df = data.copy()
```

In [ ]:

```
1 # Dropping the column which is not required " serial No"
```

In [72]:

```
1 df.drop(["Serial No."],axis = 1 , inplace = True)
```

In [73]:

```
1 # Check null values
2 df.isna().sum()
```

Out[73]:

```
GRE Score      0
TOEFL Score    0
University Rating 0
SOP            0
LOR            0
CGPA           0
Research       0
Chance of Admit 0
dtype: int64
```

In [74]:

```
1 # Information about the data type of all columns
2 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   GRE Score             500 non-null   int64
1   TOEFL Score           500 non-null   int64
2   University Rating     500 non-null   int64
3   SOP                   500 non-null   float64
4   LOR                   500 non-null   float64
5   CGPA                  500 non-null   float64
6   Research              500 non-null   int64
7   Chance of Admit       500 non-null   float64
dtypes: float64(4), int64(4)
memory usage: 31.4 KB
```

**There is no null values in any column.**

**No null value detected.**

In [75]:

```
1 # checking for the number of unique values in each columns
2 df.nunique()
```

Out[75]:

```
GRE Score      49
TOEFL Score    29
University Rating  5
SOP            9
LOR            9
CGPA          184
Research        2
Chance of Admit  61
dtype: int64
```

## Observation

**University Rating,SOP,LOR,Research are seems to be categorical variables as the number of unique values are very small.**

**rest of the features are numeric , and ordinal . (University Rating,SOP,LOR,Research are discrete ) and rest are continuous**

**also if SOP , University rating , LOR and research can be considered as numeric ordinal data.**

In [ ]:

```
1
```

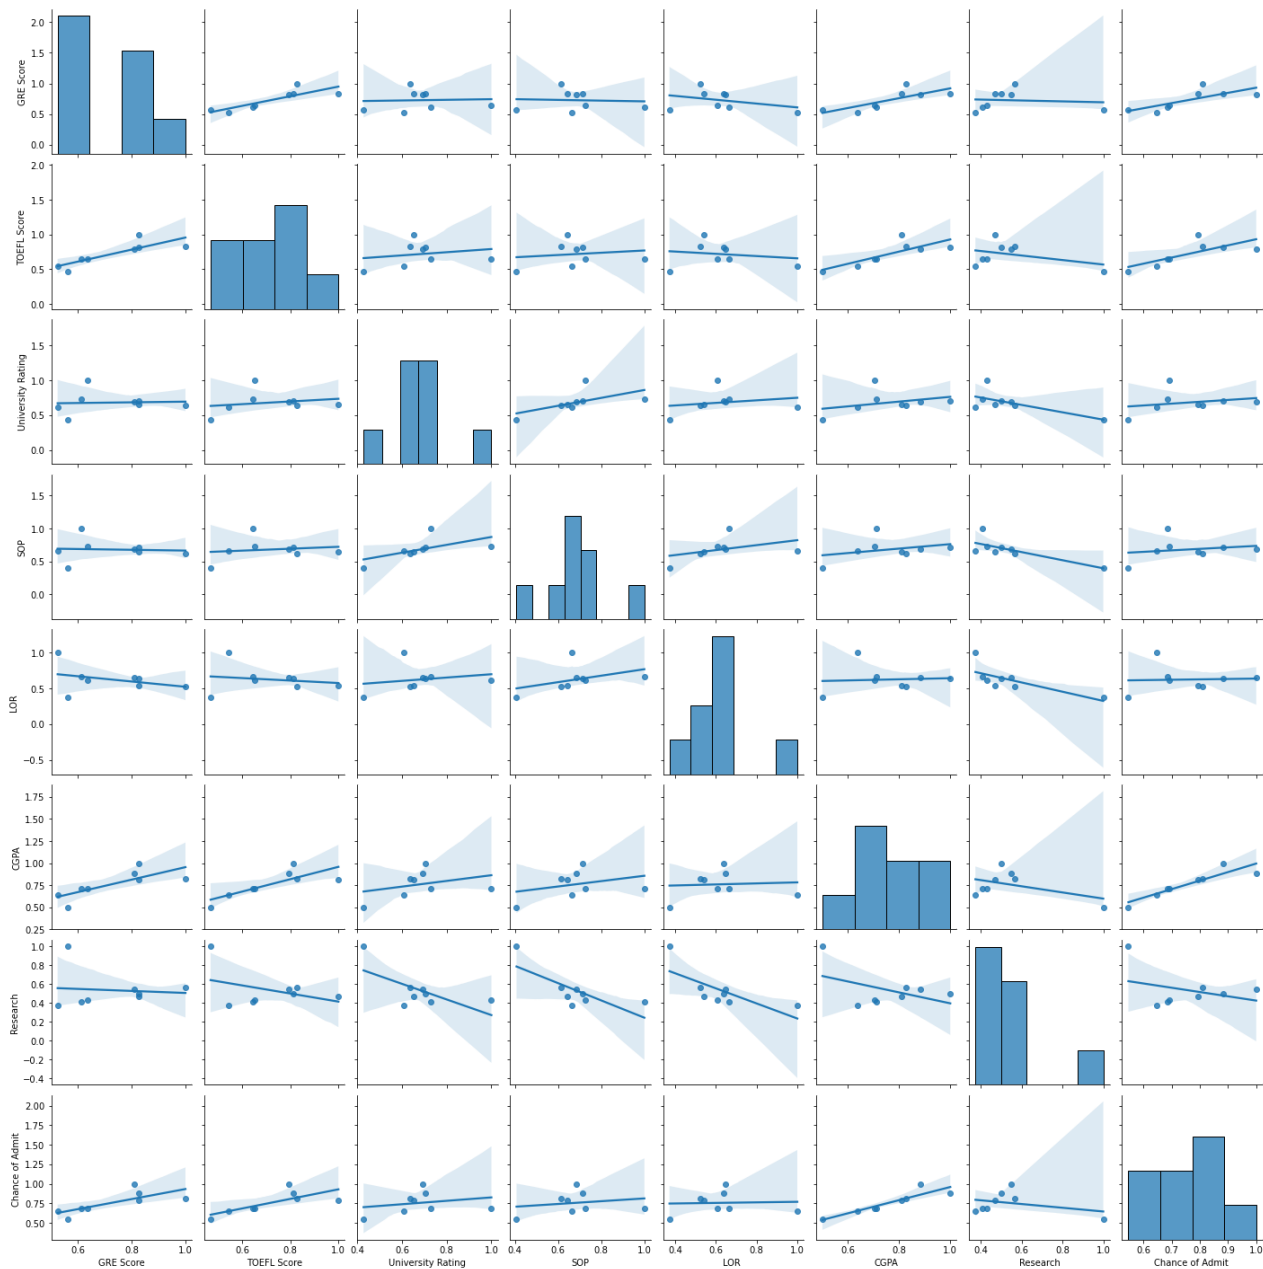
## Checking the overall linearity and correlation across all features using pairplot

In [76]:

```
1 sns.pairplot(df.corr(),kind = "reg" )
```

Out[76]:

<seaborn.axisgrid.PairGrid at 0x217e6899940>



## Overall look at the Correlation :

In [77]:


```

1 plt.figure(figsize=(9,7))
2 sns.heatmap(df.corr(),annot=True,cmap = "Greens_r")

```

Out[77]:

&lt;AxesSubplot:&gt;



## Observation

**Independent Variables (Input data):** GRE Score, TOEFL Score, University Rating, SOP, LOR, CGPA, Research

**Target/Dependent Variable :** Chance of Admit (the value we want to predict)

from above correlation heatmap , we can observe GRE score TOEFL score and CGPA have very high correlation with Change of admission.

University rating, SOP ,LOR and Research have comparatively slightly less correlated than other features.

In [ ]:

1

In [78]:

```

1 # For better convenience remove the spaces from the column names
2 df.columns

```

Out[78]:

```

Index(['GRE Score', 'TOEFL Score', 'University Rating', 'SOP', 'LOR ', 'CGPA',
      'Research', 'Chance of Admit '],
      dtype='object')

```

In [79]:

```

1 df.columns = ['GRE_Score', 'TOEFL_Score', 'University_Rating', 'SOP', 'LOR', 'CGPA',
2             'Research', 'Chance_of_Admit']

```

In [80]:

```
1 df.sample(5)
```

Out[80]:

	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Chance_of_Admit
463	304	107	3	3.5	3.0	7.86	0	0.57
158	306	106	2	2.0	2.5	8.14	0	0.61
110	305	108	5	3.0	3.0	8.48	0	0.61
288	314	104	4	5.0	5.0	9.02	0	0.82
375	304	101	2	2.0	2.5	7.66	0	0.38

In [ ]:

```
1
```

## Outliers in the Data

In [81]:

```
1 def detect_outliers(data):
2     length_before = len(data)
3     Q1 = np.percentile(data,25)
4     Q3 = np.percentile(data,75)
5     IQR = Q3-Q1
6     upperbound = Q3+1.5*IQR
7     lowerbound = Q1-1.5*IQR
8     if lowerbound < 0:
9         lowerbound = 0
10
11     length_after = len(data[(data>lowerbound)&(data<upperbound)])
12     return f"{np.round((length_before-length_after)/length_before,4)} % Outliers data from input data found"
13
14
```

In [82]:

```
1 for col in df.columns:
2     print(col, " : ",detect_outliers(df[col]))
```

```
GRE_Score : 0.0 % Outliers data from input data found
TOEFL_Score : 0.0 % Outliers data from input data found
University_Rating : 0.0 % Outliers data from input data found
SOP : 0.0 % Outliers data from input data found
LOR : 0.024 % Outliers data from input data found
CGPA : 0.0 % Outliers data from input data found
Research : 0.44 % Outliers data from input data found
Chance_of_Admit : 0.004 % Outliers data from input data found
```

In [83]:

```
1 detect_outliers(df)
```

Out[83]:

```
'0.0 % Outliers data from input data found'
```

## Observation

there are no significant amount of outliers found in the data.

In [ ]:

```
1
```

## Descriptive analysis of all numerical features :

In [84]:

```
1 df.describe()
```

Out[84]:

	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Chance_of_Admit
count	500.000000	500.000000	500.000000	500.000000	500.000000	500.000000	500.000000	500.000000
mean	316.472000	107.192000	3.114000	3.374000	3.484000	8.576440	0.560000	0.72174
std	11.295148	6.081868	1.143512	0.991004	0.92545	0.604813	0.496884	0.14114
min	290.000000	92.000000	1.000000	1.000000	1.000000	6.800000	0.000000	0.34000
25%	308.000000	103.000000	2.000000	2.500000	3.000000	8.127500	0.000000	0.63000
50%	317.000000	107.000000	3.000000	3.500000	3.500000	8.560000	1.000000	0.72000
75%	325.000000	112.000000	4.000000	4.000000	4.000000	9.040000	1.000000	0.82000
max	340.000000	120.000000	5.000000	5.000000	5.000000	9.920000	1.000000	0.97000

Observation :

chances of admit is a probability measure , which is within 0 to 1 which is good (no outliers or misleading data in column).

Range of GRE score looks like between 290 to 340.

range of TOEFL score is between 92 to 120.

university rating , SOP and LOR are distributed between range of 1 to 5.

CGPA range is between 6.8 to 9.92.

In [ ]:

```
1
```

Type *Markdown* and LaTeX:  $\alpha^2$

Graphical Analysis :

Distributions / Histogram and count plot :



In [ ]:

```
1
```

## Chance\_of\_Admit

In [17]:

```
1
2 sns.distplot(df["Chance_of_Admit"],bins = 30)
3 sm.qqplot(df["Chance_of_Admit"],fit=True, line="45")
4 plt.show()
5
```



In [ ]:



```
1
```



## GRE\_Score

In [18]:

```
1 sns.distplot(df["GRE_Score"], bins = 30)
2 sm.qqplot(df["GRE_Score"], fit=True, line="45")
3 plt.show()
```




In [ ]:

```
1
```

In [19]:

```
1 data["GRE_SCORE_CATEGORY"] = pd.qcut(data["GRE Score"], 20)
2 plt.figure(figsize=(14, 5))
3 sns.boxplot(y = data["Chance of Admit"], x = data["GRE_SCORE_CATEGORY"])
4 plt.xticks(rotation = 90)
5 plt.show()
```



Observation :

**From above boxplot (distribution of chance of admission (probability of getting admission) as per GRE score ) :**


**with higher GRE score , there is high probability of getting an admission .**

In [20]:

```
1 sns.jointplot(df["GRE_Score"],df["Chance_of_Admit"], kind = "reg",color = "brown" )  
2
```

Out[20]:

<seaborn.axisgrid.JointGrid at 0x217ecfd5df0>




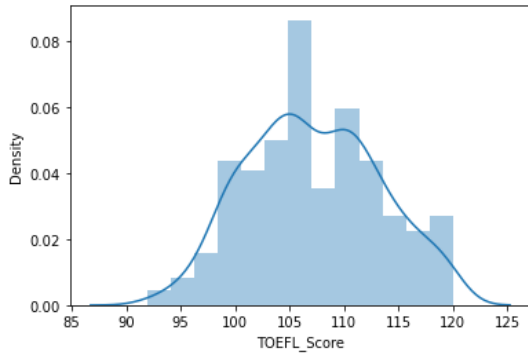
**TOEFL\_Score**

In [21]:

```


1 # TOEFL_Score
2
3 sns.distplot(df["TOEFL_Score"])
4 sm.qqplot(df["TOEFL_Score"],fit=True, line="45")
5 plt.show()
6 plt.figure(figsize=(14,5))
7 sns.boxplot(y = df["Chance_of_Admit"], x = df["TOEFL_Score"])
8

```



Out[21]:

<AxesSubplot:xlabel='TOEFL\_Score', ylabel='Chance\_of\_Admit'>

**Observation :**

**Students having high toefl score , has higher probability of getting admition .**

In [ ]:

1



In [ ]:

1

## CGPA

In [22]:

```
1  
2  
3 sns.distplot(df["CGPA"], bins = 30)  
4 sm.qqplot(df["CGPA"],fit=True, line="45")  
5 plt.show()  
6
```



## Observation :

Chance of admit and GRE score are nearly normally distributed.

GRE score, TOEFL score and CGPA has a strong correlation with chance of admission .

In [ ]:

```
1
```

In [ ]:

```
1
```

## Distribution of all other categorical features :

In [23]:

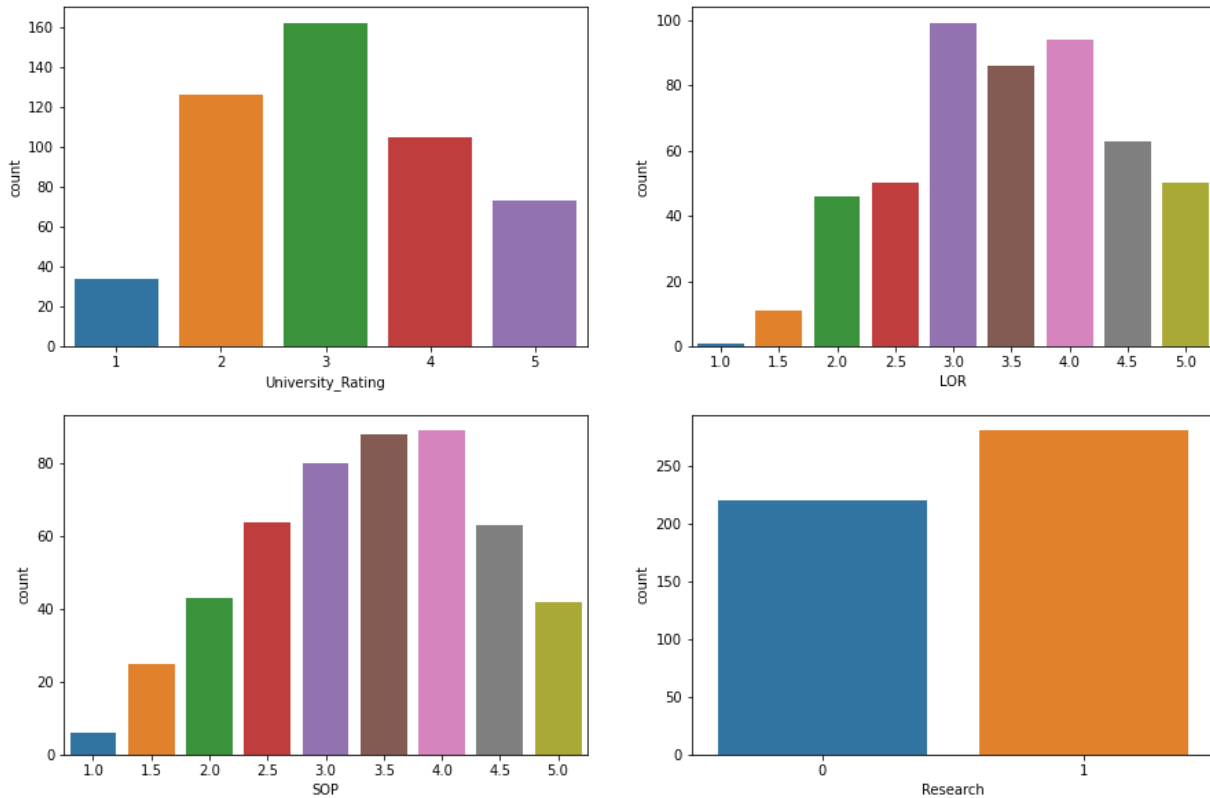
```

1 plt.figure(figsize=(15,10))
2 plt.subplot(2,2,1)
3 sns.countplot(df["University_Rating"])
4 plt.subplot(2,2,2)
5 sns.countplot(df["LOR"])
6 plt.subplot(2,2,3)
7 sns.countplot(df["SOP"])
8 plt.subplot(2,2,4)
9 sns.countplot(df["Research"])

```

Out[23]:

&lt;AxesSubplot:xlabel='Research', ylabel='count'&gt;



In [ ]:


1

In [24]:

```

1 sns.pairplot(df,y_vars = ["Chance_of_Admit"])
2 plt.title("Pair plot Chance of admit vs all the features")
3 plt.show()

```



In [ ]:

1

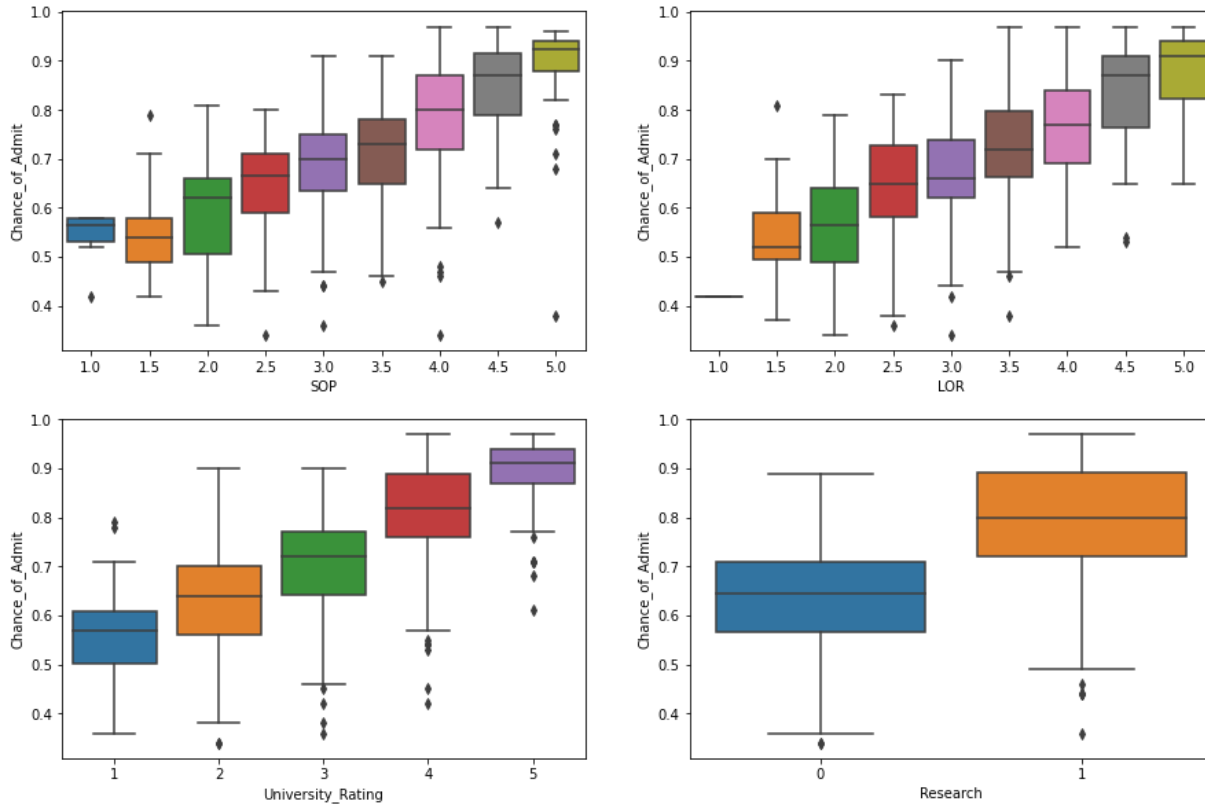
## Categorical features - vs - chances of admission boxplot :

In [25]:

```

1 plt.figure(figsize=(15,10))
2 plt.subplot(2,2,1)
3 sns.boxplot(y = df["Chance_of_Admit"], x = df["SOP"])
4 plt.subplot(2,2,2)
5 sns.boxplot(y = df["Chance_of_Admit"], x = df["LOR"])
6 plt.subplot(2,2,3)
7 sns.boxplot(y = df["Chance_of_Admit"], x = df["University_Rating"])
8 plt.subplot(2,2,4)
9 sns.boxplot(y = df["Chance_of_Admit"], x = df["Research"])
10 plt.show()

```



In [ ]:

1

## Observation:

from above plots, we can observe , statement of purpose SOP strength is positively correlated with Chance of Admission .

we can also similar pattern in Letter of Recommendation Stength and University rating , have positive correlation with Chaces of Admission .

Student having research has higher chances of Admission , but also we can observe some outliers within that caregory.

In [ ]:

1

In [ ]:

1


**Linearity : How features are correlated with Target variable - chance of admit :**

In [26]:

```
1 for col in df.columns[:-1]:
2     print(col)
3     plt.figure(figsize=(3,3))
4     sns.jointplot(df[col],df["Chance_of_Admit"],kind="reg" )
5     plt.show()
```


GRE\_Score

&lt;Figure size 216x216 with 0 Axes&gt;




TOEFL\_Score

&lt;Figure size 216x216 with 0 Axes&gt;




University\_Rating


&lt;Figure size 216x216 with 0 Axes&gt;



SOP  
<Figure size 216x216 with 0 Axes>




LOR  
<Figure size 216x216 with 0 Axes>



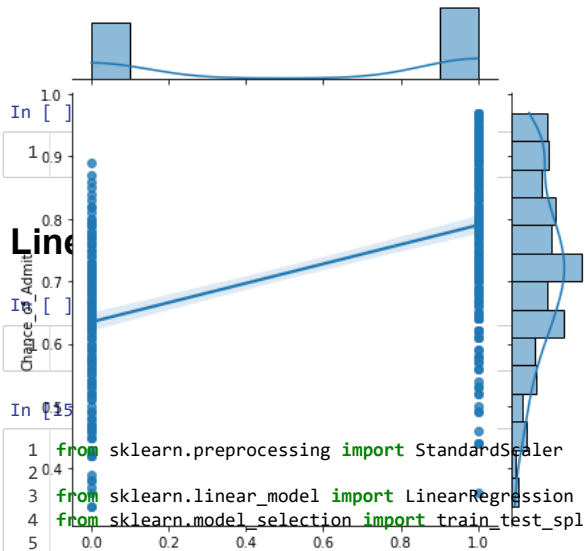
CGPA  
<Figure size 216x216 with 0 Axes>





Research

<Figure size 216x216 with 0 Axes>



```
In [15]:
1 from sklearn.preprocessing import StandardScaler
2
3 from sklearn.linear_model import LinearRegression
4 from sklearn.model_selection import train_test_split
5
6 from statsmodels.stats.outliers_influence import variance_inflation_factor
7
8 from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error, adjusted_mutual_info_score
9 from sklearn.feature_selection import f_regression
```

```
In [ ]:
1
```

```
In [ ]:
1
```

```
In [ ]:
1
```

```
In [152]:
1
2 X = df.drop(["Chance_of_Admit"],axis = 1) # independent variables
3 y = df["Chance_of_Admit"].values.reshape(-1,1) # target / dependent variables
4
5
```

```
In [ ]:
1
```

## Standardising data

In [153]:

```
1 standardizer = StandardScaler()
2 standardizer.fit(X)
3 x = standardizer.transform(X) # standardising the data
4
```

## test train splitting :

In [ ]:

```
1
```

In [154]:

```
1 X_train , X_test, y_train , y_test = train_test_split(x,y,random_state = 1,test_size = 0.2 ) # test train split
2
3
```

In [155]:

```
1 # after splitting, checking for the shape of test and train data
2
3
4 X_train.shape,X_test.shape
```

Out[155]:

```
((400, 7), (100, 7))
```

In [156]:

```
1 y_train.shape, y_test.shape
2
```

Out[156]:

```
((400, 1), (100, 1))
```

## training the model

In [157]:

```
1 LinearRegression = LinearRegression() # training LinearRegression model
2 LinearRegression.fit(X_train,y_train)
3
```

Out[157]:

```
LinearRegression()
```

## R2 score on train data :

In [158]:

```
1 r2_score(y_train,LinearRegression.predict(X_train))
```

Out[158]:

```
0.8215099192361265
```

## R2 score on test data :

In [159]:

```
1 r2_score(y_test,LinearRegression.predict(X_test) )
2
```

Out[159]:

```
0.8208741703103732
```

In [ ]:

1

## All the feature's coefficients and Intercept :

In [160]:

```
1 ws = pd.DataFrame(LinearRegression.coef_.reshape(1,-1),columns=df.columns[:-1])
2 ws["Intercept"] = LinearRegression.intercept_
3 ws
4
```

Out[160]:

	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Intercept
0	0.020675	0.019284	0.007001	0.002975	0.013338	0.070514	0.009873	0.722881

In [161]:

```
1 LinearRegression_Model_coefs = ws
2 LinearRegression_Model_coefs
3
```

Out[161]:

	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Intercept
0	0.020675	0.019284	0.007001	0.002975	0.013338	0.070514	0.009873	0.722881

In [162]:

```
1
2 def AdjustedR2score(R2,n,d):
3     return 1-(((1-R2)*(n-1))/(n-d-1))
4
```

In [163]:

```
1
2 y_pred = LinearRegression.predict(X_test)
3
4 print("MSE:",mean_squared_error(y_test,y_pred)) # MSE
5 print("RMSE:",np.sqrt(mean_squared_error(y_test,y_pred))) #RMSE
6 print("MAE :",mean_absolute_error(y_test,y_pred) ) # MAE
7 print("r2_score:",r2_score(y_test,y_pred)) # r2score
8 print("Adjusted R2 score :", AdjustedR2score(r2_score(y_test,y_pred),len(X),X.shape[1])) # adjusted R2 score
9
10
```

MSE: 0.0034590988971363833  
 RMSE: 0.058814104576507695  
 MAE : 0.040200193804157944  
 r2\_score: 0.8208741703103732  
 Adjusted R2 score : 0.8183256320830818

## Using Sklearn | Stochastic Gradient Descent Aalgorithm"

In [98]:

```
1 X = df.drop(["Chance_of_Admit"],axis = 1)
2 y = df["Chance_of_Admit"]
3 X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=2)
```

In [99]:

```
1 from sklearn.preprocessing import StandardScaler
2 scaler = StandardScaler()
3
```

In [100]:

```
1 scaler.fit(X_train)
```

Out[100]:

StandardScaler()

In [101]:

```
1 X_train = scaler.transform(X_train)
2 X_test = scaler.transform(X_test) # apply same transformation to test data
3
```

In [102]:

```
1 from sklearn.linear_model import SGDRegressor
2 from sklearn.pipeline import make_pipeline
3 sgd = make_pipeline(StandardScaler(), SGDRegressor(max_iter=1000, tol=1e-3))
```

In [103]:

```
1 sgd.fit(X_train, y_train)
2
```

Out[103]:

```
Pipeline(steps=[('standardscaler', StandardScaler()),
                  ('sgdregressor', SGDRegressor())])
```

In [104]:

```
1 y_pred = sgd.predict(X_test)
2
```

In [105]:

```
1 y_test = y_test.values
```

In [106]:

```
1 r2_score(y_test,y_pred)
```

Out[106]:

```
0.782989764614124
```

In [247]:

```
1
2 # trying different algorithms and different variations with features.
```

## Linear Regression using Statsmodel library

In [115]:

```
1 import statsmodels.api as sm
2 X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=2)
3
4
5
```

In [116]:

```
1 X_train_sm = X_train
2 X_test_sm = X_test
```

In [117]:

```
1 X_train_sm = sm.add_constant(X_train_sm)
2 X_test_sm = sm.add_constant(X_test_sm)
```

## Assumptions of linear regression

1)No multicollinearity

2)The mean of residual is nearly zero.

3)Linearity of Variables

4)Test of homoscedasticity

5)Normality of residual

In [ ]:

1

**\*Multicollinearity check :**

**checking vif scores :**

In [49]:

```
1 vifs = []
2
3 for i in range(X_train.shape[1]):
4
5     vifs.append((variance_inflation_factor(exog = X_train,
6                                             exog_idx=i)))
7 vifs
```

Out[49]:

```
[4.244635042406759,
4.063329028077076,
2.5932198746362025,
2.7053574355882417,
1.9762313259255433,
4.766976206732742,
1.466566895741921]
```

In [51]:

```
1 pd.DataFrame({ "coef_name" : X.columns ,
2               "vif" : np.around(vifs,2)})
```

Out[51]:

	coef_name :	vif :
0	GRE_Score	4.24
1	TOEFL_Score	4.06
2	University_Rating	2.59
3	SOP	2.71
4	LOR	1.98
5	CGPA	4.77
6	Research	1.47

In [118]:

```
1 olsres = sm.OLS(y_train,X_train_sm).fit()
```

In [119]:

```
1 print(olsres.summary())
```

```

              OLS Regression Results
=====
Dep. Variable:      Chance_of_Admit      R-squared:                0.829
Model:              OLS                  Adj. R-squared:           0.826
Method:              Least Squares        F-statistic:             272.1
Date:                Wed, 14 Dec 2022      Prob (F-statistic):      3.33e-146
Time:                18:27:43              Log-Likelihood:          573.41
No. Observations:    400                  AIC:                    -1131.
Df Residuals:        392                  BIC:                    -1099.
Df Model:            7
Covariance Type:     nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const          -1.3418      0.116    -11.613      0.000     -1.569     -1.115
GRE_Score       0.0021      0.001      3.893      0.000      0.001      0.003
TOEFL_Score     0.0030      0.001      3.024      0.003      0.001      0.005
University_Rating 0.0048      0.004      1.185      0.237     -0.003      0.013
SOP             0.0021      0.005      0.428      0.669     -0.008      0.012
LOR            0.0186      0.005      4.131      0.000      0.010      0.027
CGPA            0.1134      0.011     10.633      0.000      0.092      0.134
Research        0.0247      0.007      3.476      0.001      0.011      0.039
=====
Omnibus:          94.166    Durbin-Watson:           1.943
Prob(Omnibus):    0.000    Jarque-Bera (JB):        231.309
Skew:            -1.158    Prob(JB):                5.92e-51
Kurtosis:         5.918    Cond. No.                1.33e+04
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
 [2] The condition number is large, 1.33e+04. This might indicate that there are strong multicollinearity or other numerical problems.

In [120]:

```
1 r2_score(y_test,olsres.predict(X_test_sm))
```

Out[120]:

0.7927524897595928

In [ ]:

1

## Observation :

VIF score are all below 5 , doesnt seem to have very high multicollinearity.

same result of r2 value , as sklearn OLS regressor. ,

In [ ]:

1

In [ ]:

1

## Residual analysis :

In [164]:

```

1 y_predicted = LinearRegression.predict(X_train)
2 y_predicted.shape
3
4

```

Out[164]:


(400, 1)

In [165]:

```

1 residuals = (y_train - y_predicted)
2 plt.figure(figsize=(12,4))
3 plt.subplot(1,2,1)
4 sns.distplot(residuals)
5 plt.subplot(1,2,2)
6 stats.probplot(residuals.reshape(-1,), plot = plt)
7 plt.title('QQ plot for residual')
8 plt.show()

```



In [ ]:

1


## Linearity of variables

In [166]:

```

1 sns.pairplot(df,y_vars = ["Chance_of_Admit"])
2 plt.show()

```



In [ ]:

1


## Test of homoscedasticity | plotting y\_predicted and residuals

In [173]:

```

1 # Test of homoscedasticity
2 sns.scatterplot(y_predicted.reshape(-1,), residuals.reshape(-1,))
3 plt.xlabel('y_predicted')
4 plt.ylabel('Residuals')
5 plt.axhline(y=0)
6 plt.title("Y_predicted vs residuals, check of homoscedasticity")
7 plt.show()

```



In [ ]:

1

## Observation

### Homoscedasticity

from above residual plot , we can observe the varinace is not so constant .

all residuals are not evenly distributed.

In [ ]:

1

## Model Regularisation :

In [175]:

```
1 from sklearn.linear_model import Ridge # L2 regualrization
2 from sklearn.linear_model import Lasso # L1 regualrization
3 from sklearn.linear_model import ElasticNet
```

In [ ]:

1

## L2 regularization

### Ridge regression :

In [178]:

```
1 ## Hyperparameter Tuning : for appropriate lambda value :
2
3 train_R2_score = []
4 test_R2_score = []
5 lambdas = []
6 train_test_difference_Of_R2 = []
7 lambda_ = 0
8 while lambda_ <= 5:
9     lambdas.append(lambda_)
10    RidgeModel = Ridge(lambda_)
11    RidgeModel.fit(X_train,y_train)
12    trainR2 = RidgeModel.score(X_train,y_train)
13    testR2 = RidgeModel.score(X_test,y_test)
14    train_R2_score.append(trainR2)
15    test_R2_score.append(testR2)
16
17    lambda_ += 0.01
```

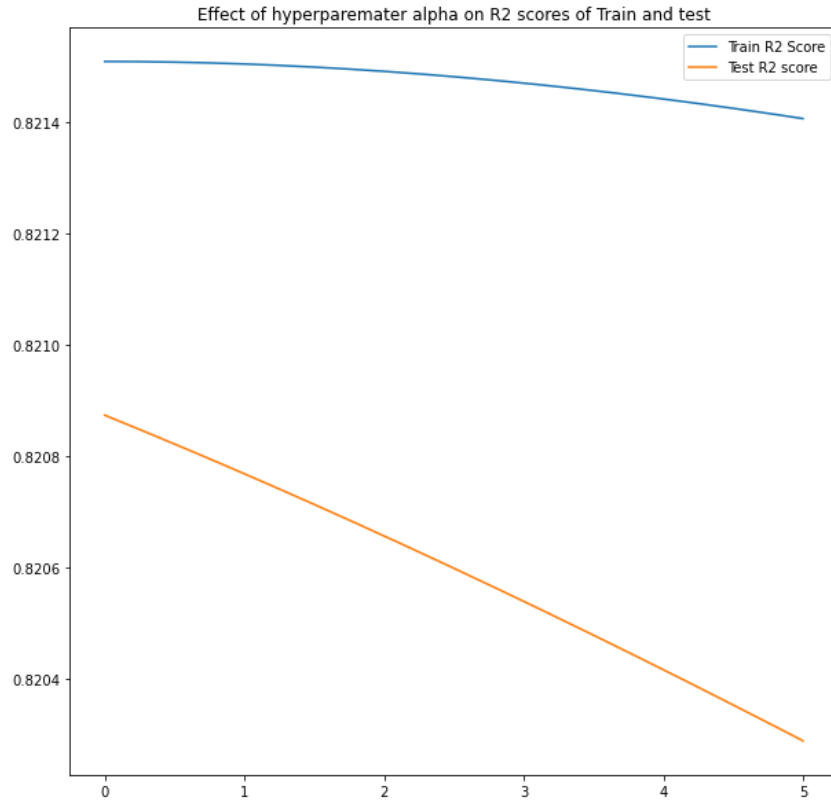


In [179]:

```

1 plt.figure(figsize = (10,10))
2 sns.lineplot(lambdas,train_R2_score,)
3 sns.lineplot(lambdas, test_R2_score)
4 plt.legend(['Train R2 Score','Test R2 score'])
5 plt.title("Effect of hyperparemater alpha on R2 scores of Train and test")
6
7
8
9 plt.show()
10

```



In [180]:

```

1 RidgeModel = Ridge(alpha = 0.1)
2 RidgeModel.fit(X_train,y_train)
3 trainR2 = RidgeModel.score(X_train,y_train)
4 testR2 = RidgeModel.score(X_test,y_test)

```

In [181]:

```

1
2 trainR2,testR2

```

Out[181]:

```
(0.8215098726041209, 0.8208639536156423)
```

In [182]:

```

1
2 RidgeModel.coef_

```

Out[182]:

```
array([[0.02069489, 0.01929637, 0.00700953, 0.00298992, 0.01334235,
        0.07044884, 0.00987467]])
```

In [183]:

```

1 RidgeModel_coefs = pd.DataFrame(RidgeModel.coef_.reshape(1,-1),columns=df.columns[:-1])
2 RidgeModel_coefs["Intercept"] = RidgeModel.intercept_
3 RidgeModel_coefs

```

Out[183]:

	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Intercept
0	0.020695	0.019296	0.00701	0.00299	0.013342	0.070449	0.009875	0.722882

In [184]:

```
1 LinearRegression_Model_coefs
```

Out[184]:

	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Intercept
0	0.020675	0.019284	0.007001	0.002975	0.013338	0.070514	0.009873	0.722881

In [185]:

```
1
2 y_pred = RidgeModel.predict(X_test)
3
4 print("MSE:",mean_squared_error(y_test,y_pred)) # MSE
5 print("RMSE:",np.sqrt(mean_squared_error(y_test,y_pred))) #RMSE
6 print("MAE :",mean_absolute_error(y_test,y_pred) ) # MAE
7 print("r2_score:",r2_score(y_test,y_pred)) # r2score
8 print("Adjusted R2 score :", AdjustedR2score(r2_score(y_test,y_pred),len(X),X.shape[1])) # adjusted R2 score
9
```


MSE: 0.003459296191728331  
 RMSE: 0.058815781825359854  
 MAE : 0.040203055117056935  
 r2\_score: 0.8208639536156423  
 Adjusted R2 score : 0.818315270028873

In [ ]:

```
1
```

In [186]:

```
1
2 y_predicted = RidgeModel.predict(X_train)
3
4 residuals = (y_train - y_predicted)
5 plt.figure(figsize=(12,4))
6 plt.subplot(1,2,1)
7 sns.distplot(residuals)
8 plt.subplot(1,2,2)
9 stats.probplot(residuals.reshape(-1,), plot = plt)
10 plt.title('QQ plot for residual')
11 plt.show()
12
```



In [ ]:

```
1
```

## L1 regularization :

## Lasso :

In [189]:


```
1  ## Hyperparameter Tuning : for appropriate lambda value :
2
3  train_R2_score = []
4  test_R2_score = []
5  lambdas = []
6  train_test_difference_Of_R2 = []
7  lambda_ = 0
8  while lambda_ <= 5:
9      lambdas.append(lambda_)
10     LassoModel = Lasso(alpha=lambda_)
11     LassoModel.fit(X_train , y_train)
12     trainR2 = LassoModel.score(X_train,y_train)
13     testR2 = LassoModel.score(X_test,y_test)
14     train_R2_score.append(trainR2)
15     test_R2_score.append(testR2)
16
17     lambda_ += 0.001
```

In [ ]:

1

In [190]:

```
1
2  plt.figure(figsize = (10,10))
3  sns.lineplot(lambdas,train_R2_score,)
4  sns.lineplot(lambdas, test_R2_score)
5  plt.legend(['Train R2 Score','Test R2 score'])
6  plt.title("Effect of hyperparemater alpha on R2 scores of Train and test")
7
8
9  plt.show()
10
```



In [ ]:

1

In [191]:

```

1 LassoModel = Lasso(alpha=0.001)
2 LassoModel.fit(X_train , y_train)
3 trainR2 = LassoModel.score(X_train,y_train)
4 testR2 = LassoModel.score(X_test,y_test)

```

In [192]:

```
1 trainR2, testR2
```

Out[192]:

(0.82142983289567, 0.8198472607571161)

In [193]:

```

1
2 Lasso_Model_coefs = pd.DataFrame(LassoModel.coef_.reshape(1,-1),columns=df.columns[:-1])
3 Lasso_Model_coefs["Intercept"] = LassoModel.intercept_
4 Lasso_Model_coefs

```

Out[193]:

	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Intercept
0	0.020616	0.019069	0.006782	0.002808	0.012903	0.070605	0.009278	0.722863

In [194]:

```
1 RidgeModel_coefs
```

Out[194]:

	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Intercept
0	0.020695	0.019296	0.00701	0.00299	0.013342	0.070449	0.009875	0.722882

In [195]:

```
1 LinearRegression_Model_coefs
```

Out[195]:

	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Intercept
0	0.020675	0.019284	0.007001	0.002975	0.013338	0.070514	0.009873	0.722881

In [ ]:


```
1
```

In [196]:

```

1
2 y_predicted = LassoModel.predict(X_train)
3
4 residuals = (y_train - y_predicted)
5 plt.figure(figsize=(12,4))
6 plt.subplot(1,2,1)
7 sns.distplot(residuals)
8 plt.subplot(1,2,2)
9 stats.probplot(residuals.reshape(-1,), plot = plt)
10 plt.title('QQ plot for residual')
11 plt.show()
12

```



In [197]:

```

1
2 y_pred = LassoModel.predict(X_test)
3
4 print("MSE:",mean_squared_error(y_test,y_pred)) # MSE
5 print("RMSE:",np.sqrt(mean_squared_error(y_test,y_pred))) #RMSE
6 print("MAE :",mean_absolute_error(y_test,y_pred) ) # MAE
7 print("r2_score:",r2_score(y_test,y_pred)) # r2score
8 print("Adjusted R2 score :", AdjustedR2score(r2_score(y_test,y_pred),len(X),X.shape[1])) # adjusted R2 score
9

```

```

MSE: 0.0034789295475193297
RMSE: 0.058982451182697807
MAE : 0.04022896061335951
r2_score: 0.8198472607571161
Adjusted R2 score : 0.8172841120280507

```

In [ ]:

1

## ElasticNet

### L1 and L2 regularisation :

Elastic net linear regression uses the penalties from both the lasso and ridge techniques to regularize regression models.

In [201]:

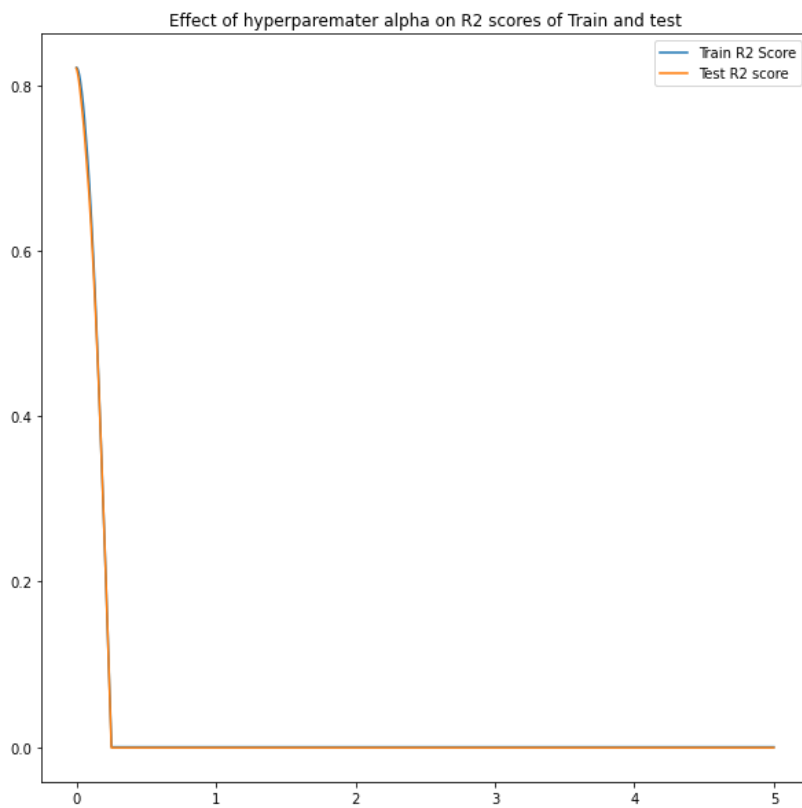
```

1
2 ## Hyperparameter Tuning : for appropriate lambda value :
3
4 train_R2_score = []
5 test_R2_score = []
6 lambdas = []
7 train_test_difference_Of_R2 = []
8 lambda_ = 0
9 while lambda_ <= 5:
10     lambdas.append(lambda_)
11     ElasticNet_model = ElasticNet(alpha=lambda_)
12     ElasticNet_model.fit(X_train , y_train)
13     trainR2 = ElasticNet_model.score(X_train,y_train)
14     testR2 = ElasticNet_model.score(X_test,y_test)
15     train_R2_score.append(trainR2)
16     test_R2_score.append(testR2)
17
18     lambda_ += 0.001
19

```

In [202]:

```
1 plt.figure(figsize = (10,10))
2 sns.lineplot(lambdas,train_R2_score,)
3 sns.lineplot(lambdas, test_R2_score)
4 plt.legend(['Train R2 Score','Test R2 score'])
5 plt.title("Effect of hyperparemater alpha on R2 scores of Train and test")
6
7
8 plt.show()
```



In [203]:

```
1 ElasticNet_model = ElasticNet(alpha=0.001)
2 ElasticNet_model.fit(X_train , y_train)
3 trainR2 = ElasticNet_model.score(X_train,y_train)
4 testR2 = ElasticNet_model.score(X_test,y_test)
```

In [204]:

```
1
2 trainR2,testR2
```

Out[204]:


```
(0.8214893364453533, 0.8203602261096284)
```

In [205]:

```

1 y_predicted = ElasticNet_model.predict(X_train)
2
3 residuals = (y_train - y_predicted)
4 plt.figure(figsize=(12,4))
5 plt.subplot(1,2,1)
6 sns.distplot(residuals)
7 plt.subplot(1,2,2)
8 stats.probplot(residuals.reshape(-1,), plot = plt)
9 plt.title('QQ plot for residual')
10 plt.show()
11

```



In [206]:

```

1
2 y_pred = ElasticNet_model.predict(X_test)
3
4 print("MSE:", mean_squared_error(y_test, y_pred)) # MSE
5 print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred))) # RMSE
6 print("MAE :", mean_absolute_error(y_test, y_pred)) # MAE
7 print("r2_score:", r2_score(y_test, y_pred)) # r2score
8 print("Adjusted R2 score :", AdjustedR2score(r2_score(y_test, y_pred), len(X), X.shape[1])) # adjusted R2 score
9

```

MSE: 0.003469023673596966  
 RMSE: 0.058898418260569324  
 MAE : 0.04021407699792928  
 r2\_score: 0.8203602261096284  
 Adjusted R2 score : 0.8178043756680987

In [207]:

```

1
2 ElasticNet_model_coefs = pd.DataFrame(ElasticNet_model.coef_.reshape(1, -1), columns=df.columns[:-1])
3 ElasticNet_model_coefs["Intercept"] = ElasticNet_model.intercept_
4 ElasticNet_model_coefs
5

```

Out[207]:

	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Intercept
0	0.020679	0.019199	0.006908	0.00292	0.013128	0.070437	0.009581	0.722873

In [208]:

```
1 RidgeModel_coefs
```

Out[208]:

	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Intercept
0	0.020695	0.019296	0.00701	0.00299	0.013342	0.070449	0.009875	0.722882

In [209]:

```
1 Lasso_Model_coefs
```

Out[209]:

	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Intercept
0	0.020616	0.019069	0.006782	0.002808	0.012903	0.070605	0.009278	0.722863

In [210]:

```
1
2 LinearRegression_Model_coefs
```

Out[210]:

	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Intercept
0	0.020675	0.019284	0.007001	0.002975	0.013338	0.070514	0.009873	0.722881

In [ ]:

```
1
```

In [211]:

```
1
2 y_pred = ElasticNet_model.predict(X_test)
3 ElasticNet_model_metrics = []
4 ElasticNet_model_metrics.append(mean_squared_error(y_test,y_pred)) # MSE
5 ElasticNet_model_metrics.append(np.sqrt(mean_squared_error(y_test,y_pred))) #RMSE
6 ElasticNet_model_metrics.append(mean_absolute_error(y_test,y_pred) ) # MAE
7 ElasticNet_model_metrics.append(r2_score(y_test,y_pred)) # r2score
8 ElasticNet_model_metrics.append(AdjustedR2score(r2_score(y_test,y_pred),len(X),X.shape[1])) # adjusted R2 score
9
```

In [ ]:

```
1
```

In [212]:

```
1 y_pred = LinearRegression.predict(X_test)
2 LinearRegression_model_metrics = []
3 LinearRegression_model_metrics.append(mean_squared_error(y_test,y_pred)) # MSE
4 LinearRegression_model_metrics.append(np.sqrt(mean_squared_error(y_test,y_pred))) #RMSE
5 LinearRegression_model_metrics.append(mean_absolute_error(y_test,y_pred) ) # MAE
6 LinearRegression_model_metrics.append(r2_score(y_test,y_pred)) # r2score
7 LinearRegression_model_metrics.append(AdjustedR2score(r2_score(y_test,y_pred),len(X),X.shape[1])) # adjusted R2 score
8
```

In [ ]:

```
1
```

In [213]:

```
1
2 y_pred = RidgeModel.predict(X_test)
3 RidgeModel_model_metrics = []
4 RidgeModel_model_metrics.append(mean_squared_error(y_test,y_pred)) # MSE
5 RidgeModel_model_metrics.append(np.sqrt(mean_squared_error(y_test,y_pred))) #RMSE
6 RidgeModel_model_metrics.append(mean_absolute_error(y_test,y_pred) ) # MAE
7 RidgeModel_model_metrics.append(r2_score(y_test,y_pred)) # r2score
8 RidgeModel_model_metrics.append(AdjustedR2score(r2_score(y_test,y_pred),len(X),X.shape[1])) # adjusted R2 score
9
```

In [ ]:

```
1
```

In [214]:

```
1
2 y_pred = LassoModel.predict(X_test)
3 LassoModel_model_metrics = []
4 LassoModel_model_metrics.append(mean_squared_error(y_test,y_pred)) # MSE
5 LassoModel_model_metrics.append(np.sqrt(mean_squared_error(y_test,y_pred))) #RMSE
6 LassoModel_model_metrics.append(mean_absolute_error(y_test,y_pred) ) # MAE
7 LassoModel_model_metrics.append(r2_score(y_test,y_pred)) # r2score
8 LassoModel_model_metrics.append(AdjustedR2score(r2_score(y_test,y_pred),len(X),X.shape[1])) # adjusted R2 score
9
```

In [ ]:

```
1
```



In [215]:

```
1 ElasticNet_model_metrics
```

Out[215]:

```
[0.003469023673596966,
0.058898418260569324,
0.04021407699792928,
0.8203602261096284,
0.8178043756680987]
```

In [ ]:

```
1
```

In [216]:

```
1
2 A = pd.DataFrame([LinearRegression_model_metrics,LassoModel_model_metrics,RidgeModel_model_metrics,ElasticNet_model_metrics],colu
3 A
```

Out[216]:

	MSE	RMSE	MAE	R2_SCORE	ADJUSTED_R2
Linear Regression Model	0.003459	0.058814	0.040200	0.820874	0.818326
Lasso Regression Model	0.003479	0.058982	0.040229	0.819847	0.817284
Ridge Regression Model	0.003459	0.058816	0.040203	0.820864	0.818315
ElasticNet Regression Model	0.003469	0.058898	0.040214	0.820360	0.817804

In [217]:

```
1 B = pd.DataFrame(LinearRegression_Model_coefs.append(Lasso_Model_coefs).append(RidgeModel_coefs).append(ElasticNet_model_coefs))
2 B.index = ["Linear Regression Model","Lasso Regression Model","Ridge Regression Model","ElasticNet Regression Model"]
3
```

In [218]:

```
1 REPORT = B.reset_index().merge(A.reset_index())
```

In [219]:

```
1
2 REPORT = REPORT.set_index("index")
3 REPORT
```

Out[219]:


	GRE_Score	TOEFL_Score	University_Rating	SOP	LOR	CGPA	Research	Intercept	MSE	RMSE	MAE	R2_SCORE
index												
Linear Regression Model	0.020675	0.019284	0.007001	0.002975	0.013338	0.070514	0.009873	0.722881	0.003459	0.058814	0.040200	0.82087
Lasso Regression Model	0.020616	0.019069	0.006782	0.002808	0.012903	0.070605	0.009278	0.722863	0.003479	0.058982	0.040229	0.81984
Ridge Regression Model	0.020695	0.019296	0.007010	0.002990	0.013342	0.070449	0.009875	0.722882	0.003459	0.058816	0.040203	0.82086
ElasticNet Regression Model	0.020679	0.019199	0.006908	0.002920	0.013128	0.070437	0.009581	0.722873	0.003469	0.058898	0.040214	0.82036

In [221]:

```
1 sns.barplot(y = REPORT.loc["ElasticNet Regression Model"][0:7].index,
2             x = REPORT.loc["ElasticNet Regression Model"][0:7])
```

Out[221]:

&lt;AxesSubplot:xlabel='ElasticNet Regression Model'&gt;



In [ ]:

1

In [ ]:

1

## Insights , Feature Importance and Interpretations and Recommendations :

first column was observed as unique row identifier which was dropped and was not required for model building.

University Rating , SOP and LOR strength and research are seems to be discrete random Variables , but also ordinal numeric data.

all the other features are numeric, ordinal and continuous.

No null values were present in data.

No Significant amount of outliers were found in data.

Chance of admission(target variable) and GRE score(an independent feature) are nearly normally distributed.

Independent Variables (Input data): GRE Score, TOEFL Score, University Rating, SOP, LOR, CGPA, Research

Target/Dependent Variable : Chance of Admit (the value we want to predict)

from correlation heatmap , we can observe GRE score, TOEFL score and CGPA have very high correlation with Change of admission.

University rating, SOP ,LOR and Research have comparatively slightly less correlated than other features.

chances of admit is a probability measure , which is within 0 to 1 which is good (no outliers or misleading data in column).

Range of GRE score looks like between 290 to 340.

range of TOEFL score is between 92 to 120.

university rating , SOP and LOR are distributed between range of 1 to 5.

CGPA range is between 6.8 to 9.92.

From boxplots (distribution of chance of admission (probability of getting admission) as per GRE score ) : with higher GRE score , there is high probability of getting an admission .

Students having high toefl score , has higher probability of getting admission .

from count plots, we can observe , statement of purpose SOP strength is positively correlated with Chance of Admission .

we can also similar pattern in Letter of Recommendation Stength and University rating , have positive correlation with Chaces of Admission .

Student having research has higher chances of Admission , but also we can observe some outliers within that category.

In [ ]:

1

## Actionable Insights and Recommendations :

education institute can not just help student to improve their CGPA score but also assist them writing good LOR and SOP thus helping them admit to better university.

The education institute can not just help student to improve their GRE Score but can also assist them writing good LOR and SOP thus helping them admit to a better University.

**Awareness of CGPA and Reserach Capabilities :** Seminars can be organised to increase the awareness regarding CGPA and Research Capabilities to enhance the chance of admit.

Any student can never change their current state of attributes so awareness and marketing campaign need to surveyed hence creating a first impression on student at undergraduate level, which wont just increase company's popularity but will also help sudent get prepared for future plans in advance.

A dashboard can be created for students whenever they logged in into your website, hence allowing a healthy competition also to create a progress report for students.

Additional features like number of hours they put in studing, watching lectures, assignments soved percentage, marks in mock test can result a better report for every student to judge themselves and improve on their own.

## Regression Analysis :

from regression analysis (above bar chart and REPORT file), we can observe the CGPA is the most Important feature for prediciing the chances of admission.

other important features are GRE and TOEFL score .

after first Regression Model, checked for Multicollinearity . Getting all the VIF scores below 5 , showing there's no high multicollinearity.

all the residuals are not perfectly normally distributed. and so residual plot we can observe some level of heteroscedasticity.

regularised model ridge and lasso both give very similar results to Linear Regression Model.

similarly ElasticNet (L1+L2) also returns very similar results. along with rest of all the model metrics.